

	id	name	dob	group_id
1	1	Jim	1975-05-07	1
2	2	Sergio	1997-03-04	1
3	3	Jackson	1996-04-06	2
10	10	Jeremy	1998-07-08	2
11	11	Jack	1994-05-06	3
15	15	Mickey	2001-02-03	-1
20	20	Alex	1993-08-09	3
30	30	Marion	1997-07-07	3
31	31	Mario	1996-03-04	4
41	41	Alexander I	1997-04-05	5
42	42	Alexander II	1985-01-01	5
	<null>	<null>	<null>	<null>
	<null>	Alexander II	<null>	<null>
			???	

```
create table if not exists person
(
    id      integer,
    name   text not null,
    dob    date not null,
    group_id integer
);
```

```
be5-public> INSERT INTO public.person (id, name, dob, group_id) VALUES (null, 'Alexander II', null, null)
[2022-10-02 09:11:06] [23502] ERROR: null value in column "dob" violates not-null constraint
[2022-10-02 09:11:06] Detail: Failing row contains (null, Alexander II, null, null).
```

preventing storing data in a wrong way  
on early stage is a big win.

	id	name	dob	group_id
1	1	Jim	1975-05-07	
2	2	Sergio	1997-03-04	
3	3	Jackson	1996-04-06	
4	10	Jeremy	1998-07-08	
5	11	Jack	1994-05-06	
6	15	Mickey	2001-02-03	
7	20	Alex	1993-08-09	
8	30	Marion	1997-07-07	
9	31	Mario	1996-03-04	
10	41	Alexander I	1997-04-05	
11	42	Alexander II	1985-01-01	
12	<null>	Alexander II	1991-08-07	

	id	name	dob	group_id
	42	Alexander II	1985-01-01	5
	42	Alexander II	1991-08-07	<null>

identifier should be unique!



?

when field is  
+ NOT NULL + UNIQUE => field can be used  
as PK

what DIFF      + NOT NULL + UNIQUE      vs PRIMARY KEY

①

Columns (4) Keys (1) Indexes (1) Foreign Keys Grants

Name: id Type: integer Default:  
 Not null  Auto inc  Unique  Primary key

name text  
dob date  
group\_id integer

```
create sequence person_id_seq;
alter table person
    alter column id set default nextval('public.person_id_seq')::regclass;
```

	id	name	dob	group_id
1	1	Jim	1975-05-07	1
2	2	Sergio	1997-03-04	1
3	3	Jackson	1996-04-06	2
4	10	Jeremy	1998-07-08	2
5	11	Jack	1994-05-06	3
6	15	Mickey	2001-02-03	-1
7	20	Alex	1993-08-09	3
8	30	Marion	1997-07-07	3
9	31	Mario	1996-03-04	4
10	41	Alexander I	1997-04-05	5
11	42	Alexander II	1985-01-01	5
12	43	Alexander II	1991-08-07	<null>
13	<default>	A3	1992-03-04	<null>

sequence starts from ①

[23505] ERROR: duplicate key value violates unique constraint "person\_pkey"  
Detail: Key (id)=(1) already exists.

43	Alexander II	1991-08-07
44	A3	1991-04-05

alter sequence person\_id\_seq restart 43;

```
INSERT INTO public.person (id, name, dob, group_id, status) VALUES (DEFAULT, 'J5', '2022-10-02', null, DEFAULT)
```

<input type="checkbox"/> id	<input type="checkbox"/> name	<input type="checkbox"/> dob	<input type="checkbox"/> group_id	<input type="checkbox"/> status
45	J5	2022-10-02	<null>	created

```
create table if not exists person
(
    id      serial primary key,
    name    text not null,
    dob     date not null,
    group_id integer,
    status   text default 'created'::text
);
```

Diagram showing the relationship between the table definition and the inserted data:

- The table definition shows the `id` column as `serial primary key`, which is highlighted with a yellow oval.
- The table definition shows the `status` column as `text default 'created'::text`, which is highlighted with a yellow oval.
- The inserted data shows the `id` value as 45, which is highlighted with a yellow oval.
- The inserted data shows the `status` value as `created`, which is highlighted with a yellow oval.
- A yellow arrow points from the table definition's `id` column to the inserted data's `id` value.
- A yellow arrow points from the table definition's `status` column to the inserted data's `status` value.

```
INSERT INTO public.person (name, dob) VALUES ('J6', '2022-11-11');
```

45 J5	2022-10-02	<null>	created
46 J6	2022-11-11	<null>	created

by serial

by default

# Foreign key

person  $\longleftrightarrow$  group

	id	name	dob	group_id
1	Jim	1975-05-07		1
2	Sergio	1997-03-04		1
3	Jackson	1996-04-06		2
10	Jeremy	1998-07-08		2
11	Jack	1994-05-06		3
15	Mickey	2001-02-03		3

	id	name
1	BE1	
2	BE2	
3	BE3	
4	BE4	
5	BE5	
6	FE1	
7	FE2	
8	FE3	

is it valid?

INFO!!

only in our brain

```
select *
from person p
join group g on g.id = p.group_id
```

our table

Table: person

Comment:

Columns (5) Keys (1) Indexes (1) Foreign Keys (1) Grants

Name: string id.

Target table: group

Update rule: Delete rule:

no action no action

Deferrability:

not deferrable

SQL Script

```
alter table person
add constraint person_group_id_fk
foreign key (group_id) references group (id);
```

field from "our-table"

are there any REQ?  $\rightarrow$  IT MUST BE PK

foreign keys 1

person\_group\_id\_fk (group\_id)  $\rightarrow$  group (id)

just name

```
create table if not exists group
(
    id serial primary key,
    name text,
    type_id integer
);
```

[23503] ERROR: insert or update on table "person" violates foreign key constraint "person\_group\_id\_fk"
Detail: Key (group\_id)=(11) is not present in table "group".

transactions

from to amount date

"Jim" "Kate" \$100.00 ~~null~~

Alex Tom ~~null~~ 13.10.22

~~customer~~ sales → ~~customer~~ → ~~address~~ → ~~street~~



use "cascade" extremely careful

bad data examples

delete \* from streets  
inserts into streets ( )...()

deleted boolean

update users

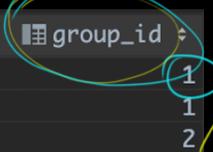
set deleted=true  
where id = 123



person

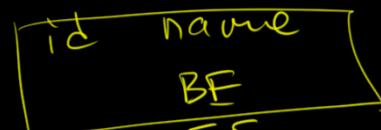
	id	name	dob	group_id
1	Jim		1975-05-07	1
2	Sergio		1997-03-04	1
3	Jackson		1996-04-06	2/null

doesn't belong to person group



1

2/null



id

name

BE

FE

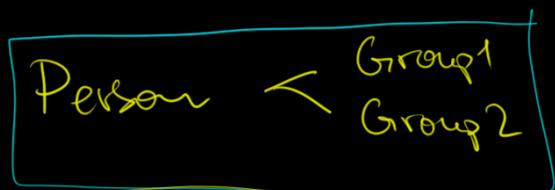
Design

Person -

Person = Group

1-1  
1:1

1-0/1



we can't  
maintain that relation  
with that structure

1-N  
1:N



	id	name	dob
1	Jim		1975-05-07
2	Sergio		1997-03-04
3	Jackson		1996-04-06

table-link

p-id	g-id
1	1
2	3
2	4
2	5

group

	id	name
1	BE1	
2	BE2	
3	BE3	
4	BE4	

each row (p-id, g-id)  
PK(p-id, g-id)

Columns (2)	Keys (1)	Indexes	Foreign
p_id int -- part of primary key g_id int -- part of primary key			

Columns (2)	Keys (1)	Indexes	Foreign
	pg_pk PRIMARY KEY (p_id, g_id)		

pg\_person\_id\_fk FOREIGN KEY (p\_id) REFERENCES person (id)  
 pg\_group\_id\_fk FOREIGN KEY (g\_id) REFERENCES group (id)

```
create table pg
(
  p_id int
    constraint pg_person_id_fk
      references person, (id)
  g_id int
    constraint pg_group_id_fk
      references group, (id)
  constraint pg_pk
    unique (p_id, g_id)
);
```

person : group  
 ↗  
 M : N  
 Many - do - Many

single responsibility

person

p_id	p.name	dob
1	Jim	1975-05-07
2	Sergio	1997-03-04
3	Jackson	1996-04-06
10	Jeremy	1998-07-08

encapsulation

pg

p_id	g_id
1	1
2	3
2	4

relation

group

g_id	g.name
1	BE1
2	BE2
3	BE3
4	BE4

encapsulation

always separate

- entities
- relations

person \* group  
 person - group  
 person < group  
 group N

group \* person  
 group - person  
 group < person  
 person N

p_id	g_id
1	1
2	3
2	4
43	4

```
select p.name
from person p
join pg on p.id = pg.p_id
join group g on g.id = pg.g_id
where g.name = 'BE4'
```

name
Sergio
Alexander II

```
select g.name
from person p
join pg on p.id = pg.p_id
join group g on g.id = pg.g_id
where p.name = 'Sergio'
```

name
BE3
BE4

ANSI SQL 92  
30 years old syntax

```
select p.name, count(p.name)
from person p
join pg on p.id = pg.p_id
join group g on g.id = pg.g_id
group by p.name
```

name	count
Sergio	2
Alexander II	1
Jim	1

```
select g.name, count(g.name)
from person p
join pg on p.id = pg.p_id
join group g on g.id = pg.g_id
group by g.name
```

name	count
BE1	1
BE3	1
BE4	2

50M records → w/o index 3-5 min  
with index 100 ms

GPS



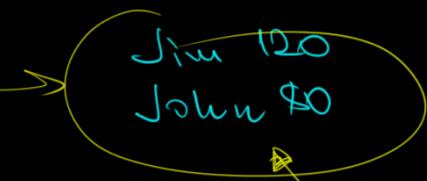
JSON

geometry  
algebra  
statistics

—

—

person-name	amount
Jim	100
John	100



we need to run them  
in transaction  
→ whether ALL  
→ whether NOTHING

company has a budget of 200

- X update salary set amount = 120 where user = 'Jim'
  - X update salary set amount = 80 where user = 'John'
- ⇒ Jim 120  
John 100
- we will not able to understand what's wrong  
immediately
- we will run out of budget.
- conn. broken here  
constraint violation  
exception

`conn.setAutoCommit(true)`

after each statement data persists into db

`conn.setAutoCommit(false);`

data will be sent to the server

But, will be persisted only after

OR, can be cancelled by

`conn.commit();`

`conn.rollback();`

You can guarantee that

set of statements are done in batch

→  
≡ }  
→ `commit();`

```
try {
    ≡ } commit()
} catch(X) {
    log...; rollback()
}
```

→ ALL  
→ NOTHING

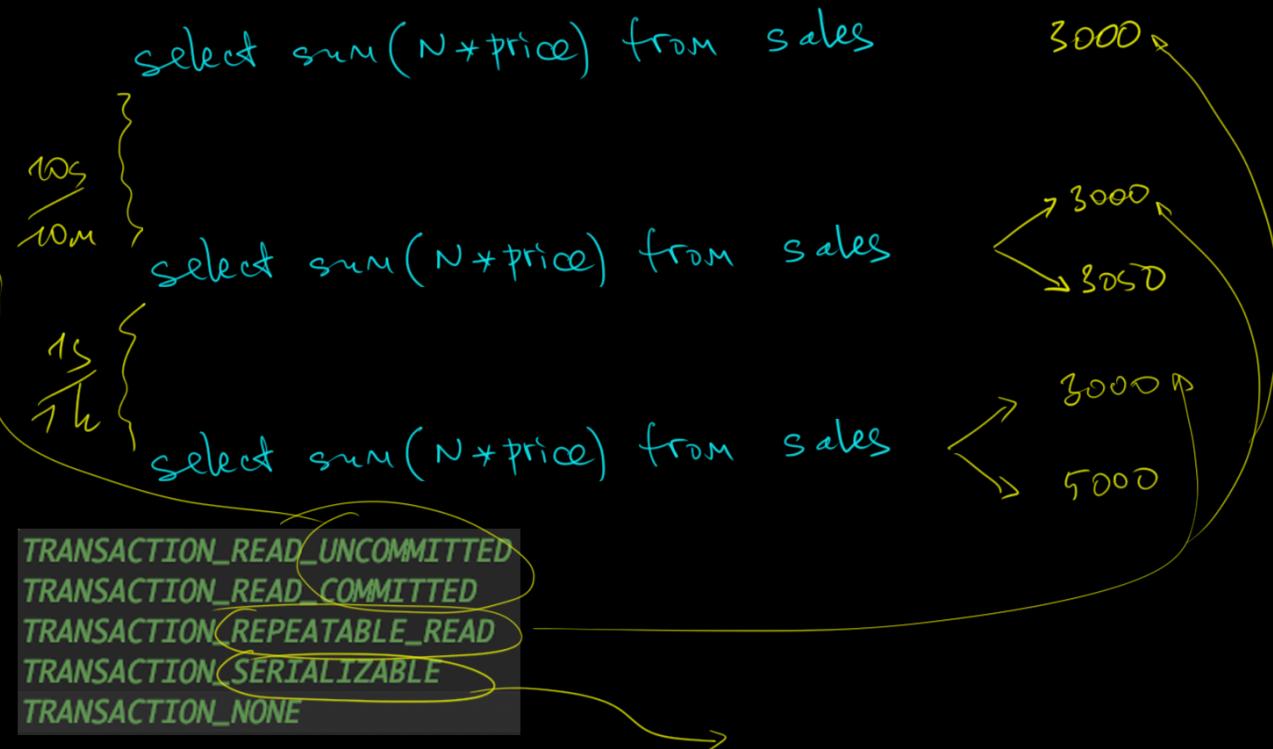
sales

	N	price
item 1	10	100
item 2	20	10 + 200
item 3	1	50
item 4	1	100

→

transaction isolation

conn.setTransactionIsolation();



select sum(N\*price) from sales < time  
 select sum(N\*price) from sales < time